




Points: estimates; lines: 95% bootstrap CI; dashed line: zero difference; Retrieval top-k label rate: $p = 0.173$; Retrieval MRR: $p = 0.036$; Readout balanced accuracy: $p = 0.32$; Learning curve at 100%: $p = 0.00293$