




Points: estimates; lines: 95% bootstrap CI; dashed line: zero difference; Retrieval top-k label rate: family-adjusted $p = 0.173$; Retrieval MRR: family-adjusted $p = 0.072$; Readout balanced accuracy: family-adjusted $p = 0.32$; Learning curve at 100%: family-adjusted $p = 0.0117$